


Single-Source Shortest Paths

Let G be a weighted directed graph where the weight attached to each edge is the cost to traverse the edge.

A path from vertex u to vertex v is a sequence of one or more edges, $\{(v_1, v_2), (v_2, v_3), \dots, (v_{r-1}, v_r)\}$, in $G.E$ where $v_1 = u$ and $v_r = v$. The cost of the path is the sum of the weights of the edges in the sequence.

The shortest path weight from u to v is the minimum cost of all paths from u to v . If there is no path from u to v then the shortest path weight is ∞ .



Shortest paths are not necessarily unique, and neither are shortest-path trees. The above figure shows a weighted, directed graph and two shortest-path trees with the same root.

Four different shortest path problems can be posed on the weighted directed graph, G .

1. From a given source vertex, s , find the shortest path weights to all other vertices in $G.V$.
2. To a given destination vertex, t , find the shortest path weights from all other vertices in $G.V$.
3. Given two vertices, u and v , find the shortest path from u to v .
4. For every pair of vertices, u and v , find the shortest path weight from u to v .

Negative weights: Suppose G has one or more edges with negative weights. Such an edge actually reduces the cost of any path using it. Suppose there's a cycle using edges whose total cost is negative. A path can run around the cycle many, many times and get any negative cost desired. If the cycle can be reached from the source vertex, s , and if the destination vertex, d , can be reached from the cycle then the shortest-path weight from s to d is $-\infty$.

INITIALIZE-SINGLE-SOURCE(G, s)

```

1  for each vertex  $v \in G.V$ 
2       $v.d = \infty$ 
3       $v.\pi = NIL$ 
4   $s.d = 0$ 



```

RELAX(u, v, w)

```

1  if  $v.d > u.d + w(u, v)$ 
2       $v.d = u.d + w(u, v)$ 
3       $v.\pi = u$ 

```




The Bellman-Ford algorithm

- Allows negative-weight edges
- Computes $v.d$ and $v.\pi$ for all $v \in V$
- Returns TRUE if no negative-weight cycles reachable from s , FALSE otherwise


BELLMAN-FORD(G, w, s)

```


1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  for  $i = 1$  to  $|G.V| - 1$ 
3      for each edge  $(u, v) \in G.E$ 
4          RELAX( $u, v, w$ )
5  for each edge  $(u, v) \in G.E$ 
6      if  $v.d > u.d + w(u, v)$ 
7          return FALSE
8  return TRUE
    
```




(a)




(b)



(c)



(d)



(e)

The nested for loops relax all edges $|V| - 1$ times. The Bellman-Ford algorithm runs in time $\Theta(|V||E|)$, since the initialization in line 1 takes $\Theta(|V|)$ time, each of the $|V| - 1$ passes over the edges in lines 2-4 takes $\Theta(|E|)$ time, and the for loop of lines 5-7 takes $\Theta(|E|)$ time.

Dijkstra's algorithm

Dijkstra's algorithm solves shortest-path problem 1 for directed weighted graph with non-negative edge weights. It is a greedy algorithm and similar to MST-Prim. Starting at the source vertex, s , it grows a tree, T , that eventually spans all vertices reachable from s . If v is a vertex in $G.V$ then $v.d$ holds the shortest-path weight from s to v . During the algorithm there are three classes of vertices: T-vertices that are in T ; A-vertices not in T but adjacent to T-vertices; and X-vertices that are not adjacent to T-vertices.

Vertices are added to T in distance order; first s , then the vertex closest to s , then the next closest, etc. If u is a T-vertex then $u.\pi$ holds the tree-parent of u and $u.d$ holds the shortest path weight from s to u .

If an A-vertex, v , is adjacent to one T-vertex, u , then $v.\pi = u$ and $v.d = u.d + w(u, v)$. If an A-vertex is adjacent to multiple T-vertices, u_1, u_2, \dots, u_r , then

$$v.d = \min(u_i.d + w(u_i, v)) \text{ for } 1 \leq i \leq r$$

to select the T-vertex, u_k , that gives the shortest s -to- v path and $v.\pi = u_k$.

If v is an X-vertex then $v.\pi = NIL$ and $v.d = \infty$.

When a vertex, u , becomes a T-vertex, $Adj[u]$ is scanned for A-vertices and X-vertices. If A-vertex, v , is adjacent to u then $v.d$ is compared to $u.d + w(u, v)$ to see if $v.d$ and $v.\pi$ need to be updated. If X-vertex, v , is adjacent to u then v is changed to an A-vertex, $v.\pi = u$, and $v.d = u.d + w(u, v)$.


A-vertices and X-vertices are stored in a priority queue, Q , ordered by their d -variables with maximum priority given to the vertex with least distance, d . As with MST-PRIM, EXTRACT-MIN(Q) extracts the vertex with least distance.

DIJKSTRA(G, w, s)


```

1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2   $S = \emptyset$ 
3   $Q = G.V$ 
4  while  $Q \neq \emptyset$ 
5       $u = \text{EXTRACT-MIN}(Q)$ 
6       $S = S \cup \{u\}$ 
7      for each vertex  $v \in G.Adj[u]$ 
8          RELAX( $u, v, w$ )


```




(a)




(b)




(c)



(d)



(e)



(f)

Like MST-PRIM, DIJKSTRA runs in $O(|E| \lg |V|)$ time.